

Random Number Generators



Algorithms typically produce numbers within $(0,1]$ or similar
Be careful with exactly zero or one

- 'Pseudo-Random' Numbers: Deterministic Algorithm
- Many Different Algorithms: linear congruential, Mersenne Twister, ...
- Built-in generators in Fortran, Python, ...
- Many others available: Gnu Scientific Library, Random123, ...
- Typically well-tested, but errors do creep in
- Be careful with initialization

Using Random Number Generators: Serial Codes

- 1 Give option to initialize random number generator
- 2 Store state of random number generator at end of run
- 3 For next run initialize with stored state
- 4 Print out state to be able to reproduce runs

Examples in Fortran and Python

- Fortran

- call `random_number(u)` produces a random number u from 0 to 1
- u can be an array
- call `random_seed(size,put,get)` sets or gets information about seed
- $size$ is an integer describing number of integers to describe state
- put is an integer array used to set the state
- get is an integer array used to get the state
- example ' call `random_seed` ' gives random initialization
- example ' call `random_seed(size=k)` '
- example ' call `random_seed(put=seed(1:k))` '

Examples in Fortran and Python

- Python
 - 'import random' sets up random number routines
 - 'randomm.random(u)' gets random numbers from $[0,1)$ (always < 1)
 - u can be an array (list)
 - 'random.seed(i)' initializes the generator with seed i
 - 'random.getstate()' gets the state of the random number generator

Example code for Python

```
import random, os, pickle

if os.path.exists('state.dat'):
    print 'Found state.dat, initializing random'
    with open('state.dat','rb') as f:
        state=pickle.load(f)
    random.setstate(state)

else:
    # Use a well-known start state
    print 'No state.dat, seeding with 1 '
    random.seed(1)

....

with open('state.dat','wb') as f:
    pickle.dump(random.getstate(),f)
```

Measurements: Integrals with Monte Carlo

Monte Carlo allows us to approximate an integral:

$$\begin{aligned} I &= \int dx F(x) \\ &= \int dx W(x) \frac{F(x)}{W(x)} \end{aligned}$$

We sample points x_i from $W(x)$ and evaluate $F(x_i)/W(x_i)$ for each point.

In the limit of a large number of samples, the average

$$\bar{X} = \frac{1}{N} \sum_{i=1, N} F(x_i)/W(x_i) \rightarrow I$$


converges to the integral I . The error typically goes like $1/\sqrt{N}$, and depends critically upon the choice of $W(x)$.

Sampling Distributions

Simple Algorithms Exist for Sampling Other Distributions:


- Different range (eg. 0 to 2π): $r = 2\pi\xi$
- Exponential Decay: $r = t_0 (-\ln \xi)$
- Linear between 0 and 1: $r = \max(\xi_1, \xi_2)$
- Quadratic between 0 and 1: $r = \max(\xi_1, \xi_2, \xi_3)$
- Gaussian (Box-Muller Algorithm):
 - ① Generates Pairs of Gaussian distributed numbers
 - ② $\Theta = 2\pi\xi_1$
 - ③ $R = \sqrt{-2\ln(\xi_2)}$
 - ④ $r_1 = R \cos(\Theta)$
 - ⑤ $r_2 = R \sin(\Theta)$

Sampling Other Distributions



- Sampling within a sphere (Rejection) :
 - ① $r_1 = 2 * (\xi_1 - 0.5)$
 - ② $r_2 = 2 * (\xi_2 - 0.5)$
 - ③ $r_3 = 2 * (\xi_3 - 0.5)$
 - ④ If $r_1^2 + r_2^2 + r_3^2 > 1$; go to (1) and repeat
- Sampling on a sphere:
 - ① Sample within a sphere
 - ② Scale x, y, and z components to be on a sphere
- Sample 3 (right-handed) orthogonal axes randomly
 - ① Sample on a sphere (set as first axis): v_1
 - ② Sample a second point within a sphere: v_a
 - ③ Set second axis: $v_2 = v_a \times v_1 / |v_a \times v_1|$
 - ④ Third axis $v_3 = v_1 \times v_2$

Sampling By Rejection



If function you want to sample has a maximum value P_{max} :

- 1 Sample x randomly in the volume
- 2 Compute the ratio of the probability $R = P(x)/P_{max}$
- 3 Select a random number, Reject if $\xi > R$, return to step 1
- 4 If $\xi < R$ you are done, keep x as the sample.

This can be extremely inefficient in large spaces with a wide dynamic range, but it can work well for low dimensions.

Sampling a general 1-D positive distribution

A general normalizable positive function D can be sampled either analytically (if closed integrals exist) or numerically.

Here we assume a positive probability D from zero to infinity:

- Let $F(X) = \int_0^X dY D(Y)$
- Note: $F(0) = 0$ and $F(\infty) = 1$
- Set $r = F^{-1}(\xi)$

If integrals or inverse is not simple, this can be done numerically.

Example: $D(x) = (1/\gamma) \exp(-x/\gamma), x > 0$

$$F(X) = 1 - \exp(-X/\gamma)$$


$$F^{-1}(\xi) = -\gamma \ln(1 - \xi) = -\gamma \ln(\xi')$$

Example: sampling a general distribution


Suppose we want to sample a momentum p from:

$$P(p) \propto p^2 \exp[-\sqrt{(\hbar c p)^2 + m^2} \Delta\tau]$$

with $m = 935 \text{ MeV}$, $\Delta\tau = 0.01 \text{ MeV}^{-1}$




$P(p)$




Integrated $P(p)$

Central Limit Theorem


Sampling a distribution from: $x_i = 100 + 15 * (\xi - 0.5)$
 ξ is uniformly distributed from 0 - 1.



100 samples



1000 samples




1,000,000 samples

(see python code histogram0.py)


Central Limit Theorem

- Assume independent samples
- Compute averages by 'blocks':
- $\bar{X}_i = \sum_{j, \dots, j+nb} X_i / nb$
- Distribution of averages over blocks converges to a Gaussian (for large enough nb)

Histogram of 100,000 blocks, each the average over 100 samples:



Central Limit Theorem



Statistical Errors from Gaussian Distributions

- Assume independent samples with gaussian distribution of 'measurements'
- Compute average $\bar{X} = \frac{1}{nb} \sum X_i$
- Compute standard error $\sigma = ((\sum (X_i - \bar{X})^2 / nb) / nb)^{1/2}$

Probabilities exact average \bar{x} of being more than 1,2,3 σ from exact integral

- $\Pr(\bar{X} - \sigma \leq x \leq \bar{X} + \sigma) \approx 0.6827$
- $\Pr(\bar{X} - 2\sigma \leq x \leq \bar{X} + 2\sigma) \approx 0.9545$
- $\Pr(\bar{X} - 3\sigma \leq x \leq \bar{X} + 3\sigma) \approx 0.9973$

Many other methods are available to try to deal with non-gaussian statistics.

Bayesian methods are valuable, particularly in difficult cases.

Classical Monte Carlo

Suppose we have N particles in a finite volume, with a Hamiltonian:

$$H = \sum_i \mathbf{p}_i^2 / (2m) + \sum_i V^1(\mathbf{x}_i) + \sum_{i < j} V^2(\mathbf{x}_i, \mathbf{x}_j)$$

We want to determine the system's properties at a finite temperature T .

The partition function is defined as the sum over all states:

$$\mathcal{Z} = \sum_i \exp [-E_i / (k_B T)]$$

From the partition function we can calculate all observables of the system in statistical equilibrium.

The states of a classical system are defined by the positions \mathbf{x}_i and the momenta \mathbf{p}_i of the particles. Each state has a definite energy E_i .

Classical Monte Carlo (cont'd)


We can sample states from the partition function using MC methods. The expectation value of an observable is

$$\langle O \rangle = \sum_i \exp [-E_i / (k_B T)] O_i$$

This can be calculated by Monte Carlo up to fairly large systems.

- Sample the states $(\mathbf{x}_i, \mathbf{p}_i)$ with energy E_i from \mathcal{Z} .
- Note that for a classical system positions and momenta can be sampled independently: for each \mathbf{x}_i the momenta can be sampled from a gaussian.
- Compute observables by averaging over these states.


Example of Classical Monte Carlo



- Create random potential (blue) as before
- Exact distribution at $\beta = 1/(k_B T) = 1$ is magenta line
- Exact distribution at $\beta = 1/(k_B T) = 5$ is cyan line


Example of Classical Monte Carlo

Sample by rejection:



100 samples per bin

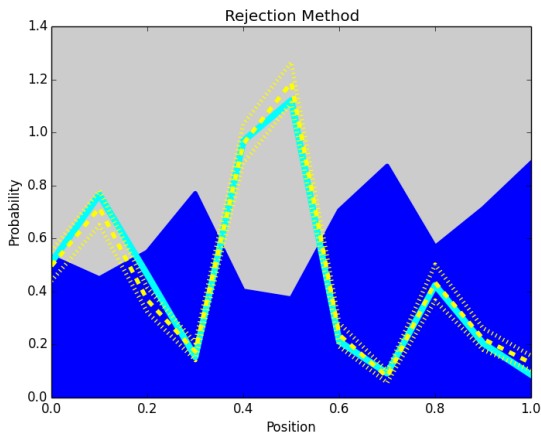
Sample by rejection for $\beta = 5$ using $\max \exp[-\beta E_i] < 1$



10000 samples

Example of Classical Monte Carlo

Sample by rejection (with errors):



Calculate averages and errors with 20 blocks, 50 samples per block
see code rejection.py

Example Code in Python

```
beta = 5 ; nsamp=50 ; nblk = 20
zpart = np.exp( -beta*y)
nexp = np.zeros ((11,nblk))
nr = np.zeros (11) ; ne = np.zeros (11)
# loop over blocks
for j in range (0,nblk):
    zr = np.random.rand(11,nsamp)
    # loop over positions for sampling
    for i in range (0,11):
        nr[i] = np.sum( zr[i,:] < zpart[i])
    # noremalize and store results for this block
    nrsum=np.sum(nr)
    nr[:] = nr[:] * 0.5 / (nrsum * 0.1)
    nexp[:,j] = nr[:]
# compute average and error for each position
for i in range(0,11):
    nr[i] = np.sum( nexp[i,:]) / nblk
    ne[i] = np.sum( (nexp[i,:]-nr[i])**2 ) / nblk
    ne[i] = np.sqrt( ne[i] / nblk )
```

Metropolis Monte Carlo and Markov Chain Algorithms

Metropolis Monte Carlo is an algorithm designed to sample complicated many-variable distributions. It employs a Markov Chain Monte Carlo algorithm to achieve this. see: *Metropolis, N.,*

Rosenbluth, A., Rosenbluth, M., Teller, A., and Teller, E., "Equations of state calculations by fast computing machines", J. Chem. Phys. 21(6) 1087 (1953).

- A Markov Chain Monte Carlo Algorithm employs random walks where each step depends only upon the present position of the system (no 'history').
- The Metropolis algorithm assumes that you want to sample from a non-negative function $W(\mathbf{R})$ which may be in many dimensions: eg. $\mathbf{R} = (\mathbf{r}_i)$ for many particles i .

Metropolis Monte Carlo

- Metropolis algorithm converges to sampling $W(\mathbf{R})$ asymptotically, but with *a priori* unknown equilibration time and unknown time between 'independent' samples.
- Detailed balance is enforced, requiring the flux from point $\mathbf{A} \rightarrow \mathbf{B}$ equals that from $\mathbf{B} \rightarrow \mathbf{A}$
- Rejection is used, we propose a move with a probability $T(\mathbf{A} \rightarrow \mathbf{B})$ and accept that move with probability $P(\mathbf{A} \rightarrow \mathbf{B})$.
- Detailed balance requires:

$$W(\mathbf{A}) T(\mathbf{A} \rightarrow \mathbf{B}) P(\mathbf{A} \rightarrow \mathbf{B}) = W(\mathbf{B}) T(\mathbf{B} \rightarrow \mathbf{A}) P(\mathbf{B} \rightarrow \mathbf{A})$$

Note that detailed balance is more restrictive than strictly necessary.

Metropolis Monte Carlo

A simple Metropolis Algorithm:

- 1 Initialize all particles within the physical volume and calculate $W(\mathbf{R})$
- 2 Propose a move of all particles \mathbf{R}' , within a box centered on the current positions: $r'_{i,j} = r_{i,j} + b(\xi_{i,j} - 0.5)$
where i runs over 3 dimensions and j over N particles
- 3 Calculate $W(\mathbf{R}')$
- 4 Accept this proposed move with probability P given by:
 $P = \min[W(\mathbf{R}')/W(\mathbf{R}), 1]$
- 5 If the move is 'accepted', set the current position to \mathbf{R}' , otherwise keep the current position at \mathbf{R}
- 6 Return to step 2, and propose a new move

After some number of steps the samples will be independent. Averages and Errors can be calculated using these independent samples

Problem for afternoon

Consider a system of 10 particles in one dimension ($0 < x_i < 1$) with a one-body (background) potential : $V_1(x) = 1 + \text{Cos}(20\pi x)$.

- Calculate analytically the probability density $\rho(x)$ using the partition function at $\beta = 0.2$ and $\beta = 1$.
- Use rejection to determine this probability density with averages and errors
- Use Metropolis Monte Carlo to sample this same density
- Add a two-body repulsive potential

$$\begin{aligned} V_{ij} &= V_0 \text{ if } |x_i - x_j| < 0.05 \\ &= 0 \text{ if } |x_i - x_j| \geq 0.05 \end{aligned}$$

Rejection is harder for the last case, what behavior do you expect?
Try $V_0 = 1$ to start with.